

Neural Networks

Monday, October 25, 2021 8:23 PM

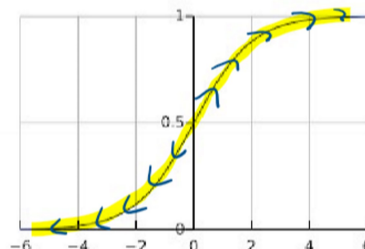
2. What is Deep Learning

- Deep learning = ML using large/deep NN
- ML = technique to create new funcs using example behavior rather than explicit instructions
- Learning from examples
 - easier to solve rather than using explicit instructions
 - examples = data, expected answer = labelled data
- ML algorithms allow to create approx functions using set of example data
- Why NN?
 - we know how to train them efficiently
 - backpropagation quickly and efficiently finds high quality approximate function
- Why Deep NN?
 - effective when we have very large data training set
 - For problems where input is unstructured, problems with complex relationships but clear goals
- Terminology
 - Data Science: process of using data analysis to build understanding
 - ML: Process of using example data to create approximate functions to be applied to new data
 - NN: ML using an interconnected network of trainable artificial neurons that map some input X to output Y
 - DL: ML using multi layered neural networks, trained w large data sets
 - Supervised learning: ML when example data provides both expected input and output, supervised by identifying and correcting mistakes
 - Labelled Data: example data with expected output, used in supervised learning
 - Unsupervised learning: ML when only expected input is provided, ML learns relationships between inputs themselves
 - Unlabeled data: example data without expected outputs, used in unsupervised learning
 - Reinforcement learning: ML using high-level goals, trial and error during training

3. ML with Logistic Regression

- Logistic Regression
 - assumes that we can make prediction (hypothesis) based on a linear combination of the outputs
- $$z = w_0x_0 + w_1x_1 + \dots + w_nx_n + b$$
- Binary Classification
 - can use 0 and 1 to represent each group
 - **Sigmoid Function**: impose function that only outputs values bw 0 and 1

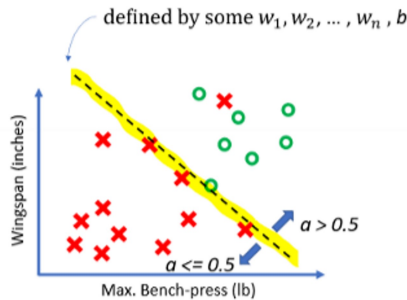
$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



- This sigmoid "forces" large values to be '1', and small values to be '0'

- Working logistic regression equation:

$$a = \sigma(w_1x_1 + w_2x_2 + \dots + w_nx_n + b)$$



- Finding parameters w_1, w_2, \dots, b

- Gradient Descent

• Cost Function J

- Compare combinations of w_n and b to know which works best
- create cost func J , measure of fitness
- if $J(w_1', w_2', \dots, b) < J(w_1, w_2, \dots, b)$ then w_1', w_2', \dots, b = better set of parameters
- Choose a cost function using accuracy
 - $Accuracy = (right\ answers / total\ answers)$
 - want a cost function that gives guidance to next guess
 - derivative of func = rate of change, where func is going
 - Gradient Descent**, adjusting parameters with understanding of where you are going

if, $J(w_1, w_2, \dots, w_n, b)$ is the **overall cost**, then

$\frac{\partial J(w_1, w_2, \dots, w_n, b)}{\partial w_1}$ is the **rate of change** of the **cost w.r.t w_1**

- can improve parameters by incrementally adjusting them based on derivative

$$w = w - \alpha \frac{\partial J(w, b)}{\partial w}$$

$$b = b - \alpha \frac{\partial J(w, b)}{\partial b}$$

Where, α , is **size of the adjustment**, normally called the **learning rate**.

- Log Function**

$$\frac{d}{da} \log(a) = \frac{1}{a}$$

- Cost for data point (Loss, L) using log func

When $y = 1$: $L(a, y) = -\log(a)$

When $y = 0$: $L(a, y) = -\log(1 - a)$

$$L(a, y) = -(y \log(a) + (1 - y) \log(1 - a))$$

- Calculate partial derivatives

$$L(a, y) = -(y \log a + (1 - y) \log(1 - a)) \Rightarrow \frac{\partial L}{\partial a} = -\frac{y}{a} + \frac{1 - y}{1 - a} = \frac{a - y}{a(1 - a)}$$

$$a = \sigma(z) = \frac{1}{1 + e^{-z}} \Rightarrow \frac{\partial a}{\partial z} = \sigma(z)(1 - \sigma(z)) = a(1 - a)$$

$$z = w_1x_1 + w_2x_2 + \dots + w_nx_n + b \Rightarrow \frac{\partial z}{\partial w_1} = x_1, \frac{\partial z}{\partial w_2} = x_2, \dots, \frac{\partial z}{\partial b} = 1$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial a} \cdot \frac{\partial a}{\partial z} \cdot \frac{\partial z}{\partial w_1} = x_1(a - y)$$

$$\frac{\partial L}{\partial w_n} = x_n(a - y)$$

$$\frac{\partial L}{\partial b} = (a - y)$$

- across each of the m data points for cost function J :

$$\frac{1}{m} \sum_{i=1}^m \dots$$

$$J = -\frac{1}{m} \left(\sum_{i=1}^m y^i \log(a^{(i)}) + \sum_{i=1}^m (1 - y^i) \log(1 - a^{(i)}) \right)$$

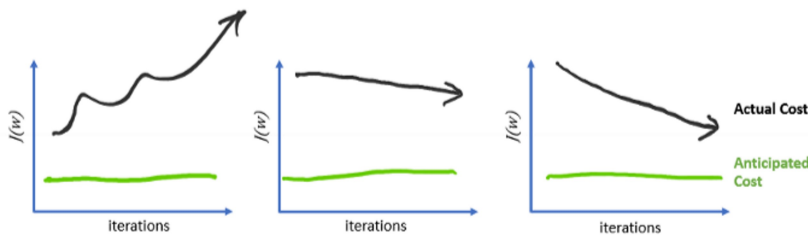
$$\frac{\partial J}{\partial w_n} = -\frac{1}{m} \sum_{i=1}^m x_n^{(i)} (a^{(i)} - y^i) \quad \frac{\partial J}{\partial b} = -\frac{1}{m} \sum_{i=1}^m (a^{(i)} - y^i)$$

• Steps:

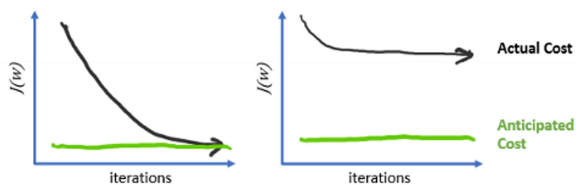
1. Assume: $a = \sigma(w_1 x_1 + w_2 x_2 + \dots + w_n x_n + b)$
2. Initialize w_{1-n}, b to random values (or zero)
3. Repeatedly apply: $w = w - \alpha \frac{\partial J(w, b)}{\partial w}$ $b = b - \alpha \frac{\partial J(w, b)}{\partial b}$
4. Stop when $J < \text{target error}$

4. Implementation and Evaluation of Logistic Regression

- Inaccuracy in ML
 - AI does not match nature of data (not linearly separable)
 - algorithm did not find best parameters
 - example data is not representative of new data
 - Not enough data to represent function
 - data is noisy
 - underlying behavior not deterministic
- Failure to find best parameters
 - accuracy impacted by ability to get to good parameters
 - **hyperparameters**
 - in logistic regression, 2 tuneable components: learning rate and number of iterations
 - Learning rate is too large
 - final parameters are worst than random
 - Learning rate is too small
 - final parameters are better than random but not optimal
 - Number of iterations is too small
 - Final parameters are better than random, but not optimal
 - Number of iteration is too large
 - As long as learning rate small enough, only costs CPU cycles



Learning rate is too **high**. Learning rate is too **low**. Number of iterations too **low**.

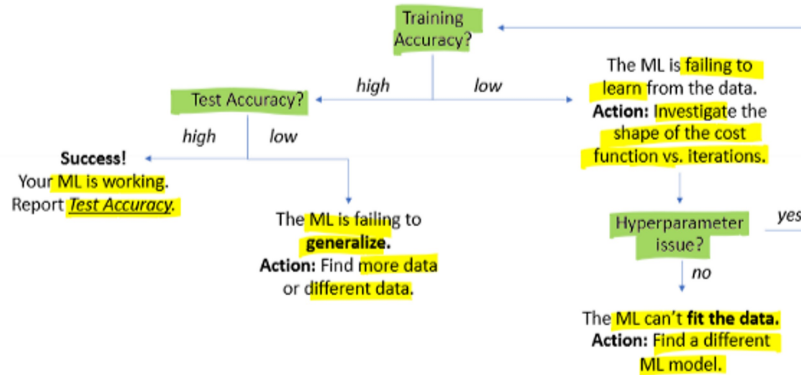


Good behavior. **Learning unsuccessful.**
(Data/model fit issue)

Issue	Symptom	Action
AI Model Doesn't Fit Data.	Training accuracy is low and hyperparameter tuning doesn't help.	Consider a different AI model (NN, for instance)
We are not finding the best parameters.	Unexpected shape of cost/iterations graph.	Tune the hyperparameters.
Example data does not represent the new data. (Lack of data, noisy data, non-deterministic data)	High training accuracy but test accuracy is low.	Try to find more data, better data or different data.

- **Building a Test Data set**

- help understand ML
- take some data, put off to the side, labelled data has the right answer



- **Vectorization**

- structure learning code correctly

$$A = \sigma(w^T X + b)$$

Handwritten annotations:
 X → inputs
 b → bias
 $w^T X + b$ → vector
 A → Answer For all inputs

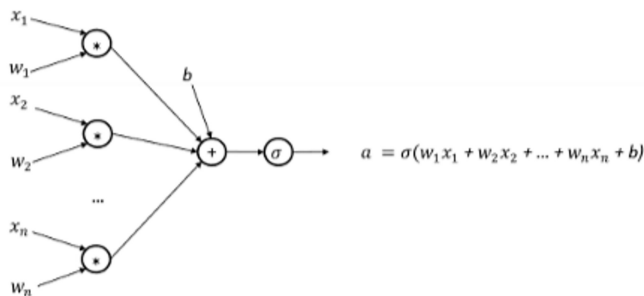
• Where:

X is a (n, m) input matrix (i.e. all of the input features for each example)
 n is the number of features in the input data
 m is the number of examples in the training data
 w is a $(n, 1)$ parameter matrix
 b is the bias
 A is a $(1, m)$ vector which represents the hypothesis for each of the m samples

5. Neural Network

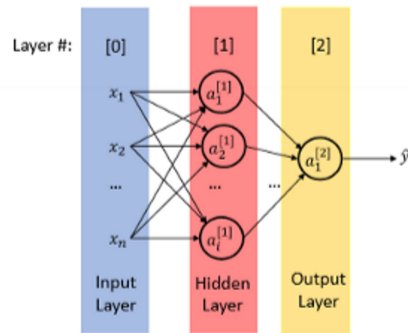
- Logistic Regression assumes a linear relationship
- NNs can learn very complex non-linear relationships rather than having to specify them

- **Computation Graph**



- **Activation Function**

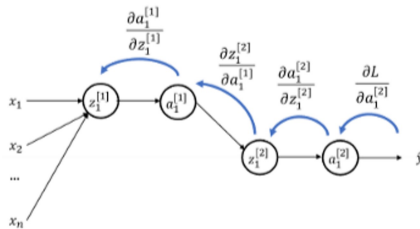
- non-linear activation function is the key to allowing combinations of logistic regression units to produce complex functions
- without non-linear activation functions, all combs of logistic regression would continue to be linear



- can learn parameters of NN using gradient descent for Logistic regression, using cost function

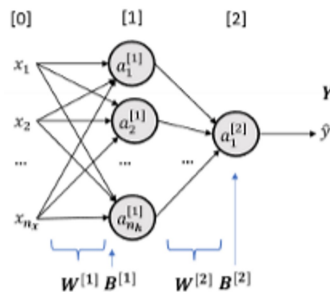
$$J(\hat{y}, y) = -\frac{1}{m} \left(\sum_{i=1}^m y^{(i)} \log(\hat{y}^{(i)}) + \sum_{i=1}^m (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right)$$

- Calculate \hat{y} using computation graph
 - determine loss
 - update each parameter (using partial derivative of cost)
 - repeat until $J < \text{target}$
- Partial derivatives and Backpropagation



- Backpropagatin the error and attributing it to each node

6. Implementation and Evaluation of NN



- If we consider a single example, i :

$$z^{[1](i)} = W^{[1]}x^{(i)} + B^{[1]}$$

$$a^{[1](i)} = g(z^{[1](i)}), \text{ where } g(z) = \tanh(z)$$

- And:

$$z^{[2](i)} = W^{[2]}a^{[1](i)} + B^{[2]}$$

$$\hat{y}^{(i)} = a^{[2](i)} = \sigma(z^{[2](i)})$$

- Backpropagation - Vectorization

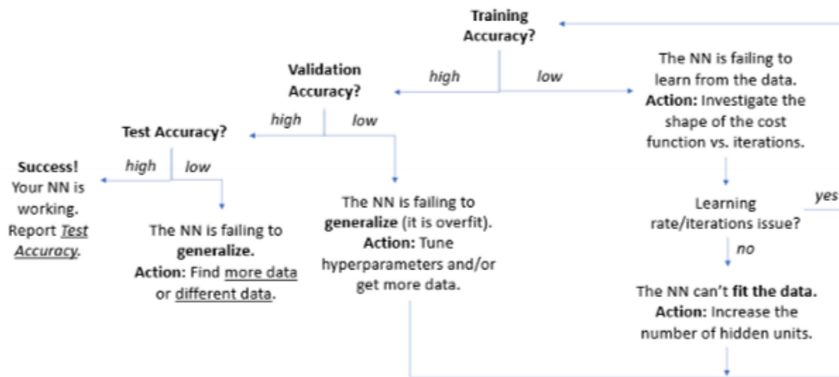
Element-wise	Vectorized - single sample, i	Vectorize for all m samples
(1) $dz_1^{[2]} = (\hat{y} - y)$	$dz^{[2](i)} = (\hat{y}^{(i)} - y^{(i)})$	$dZ^{[2]} = (\hat{Y} - Y)$ $(1, m) \quad (1, m) \quad (1, m)$
(2) $dw_{1,1}^{[2]} = dz_1^{[2]} a_1^{[1]}$	$dW^{[2](i)} = dz^{[2](i)} a^{[1](i)T}$	$dW^{[2]} = \frac{1}{m} dZ^{[2]} A^{[1]T}$ $(1, n_h) \quad (1, m) \quad (n_h, n_h)$
(3) $db_1^{[2]} = dz_1^{[2]}$	$dB^{[2](i)} = dz^{[2](i)}$	$dB^{[2]} = \frac{1}{m} \sum_{\text{rows}} dZ^{[2]}$ $(1, 1) \quad (1, m) \quad (1, m)$
(4) $dz_1^{[1]} = w_{1,1}^{[2]} dz_1^{[2]} g'(z_1^{[1]})$	$dz^{[1](i)} = W^{[2](i)T} dz^{[2](i)} + g'(z^{[1](i)})$	$dZ^{[1]} = W^{[2]T} dZ^{[2]} + g'(Z^{[1]})$ $(n_h, m) \quad (n_h, m) \quad (1, m)$
(5) $dw_{1,1}^{[1]} = dz_1^{[1]} x_1$	$dW^{[1](i)} = dz^{[1](i)} x^{(i)T}$	$dW^{[1]} = \frac{1}{m} dZ^{[1]} X^T$ $(n_h, n_x) \quad (n_h, m) \quad (m, n_x)$
(6) $db_1^{[1]} = dz_1^{[1]}$	$dB^{[1](i)} = dz^{[1](i)}$	$dB^{[1]} = \frac{1}{m} \sum_{\text{rows}} dZ^{[1]}$ $(n_h, 1) \quad (n_h, m) \quad (n_h, m)$

- Initializing parameters in NN
 - in LR, initialize parameters to any starting values, usually 0
 - Using 0 or uniform non-zero value does not work with NN?
- Overfit in NN

- it is possible that the NN finds an approx func that fits training data but not new data
- need to tune hyperparameters for optimal performance
- because of overfit in NN, split 3 data sets

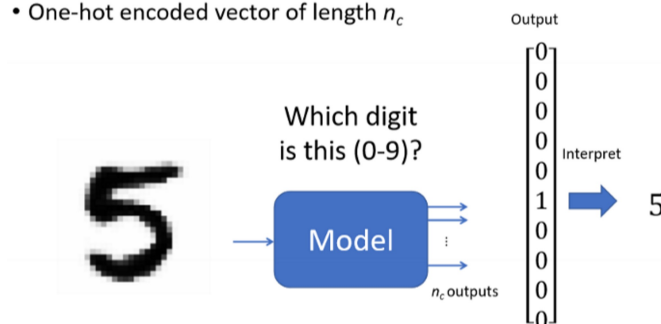


- need 3rd data set to measure expected performance
- don't want to simply fit hyperparameters to validation set
- Validation set gives data to tune hyperparameters
- Test data gives independent ref to measure performance of AI



7. Multiclass Classification

- Binary Classification
 - 2 choices
- Multiclass classification
 - n_c number of possible classes, ex. MNIST dataset
 - input has 1 label
- Multilabel classification
 - input has 1+ label
 - cannot use softmax
 - user separate classifiers or sigmoids on outputs
 - labels cannot be one hot encoded vectors
- One hot encoding
 - Binary classification model extension
 - One-hot encoded vector of length n_c

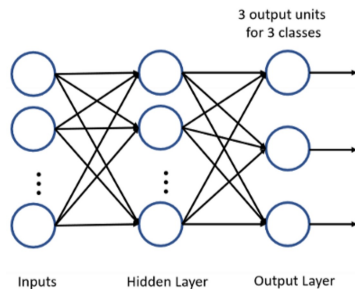


- Extending binary classifier for n_c classes
 - **Multiple Binary Classifiers**
 - One vs All
 - build multiple binary classifiers, one binary classifier/class each classifier predicts whether input is in class or not
 - One vs One

- $n_c(n_c - 1)/2$ binary classifiers, all possible combinations of 2 classes
- each classifier only receives data about pair of classes discriminating between
- majority voting scheme to select predicted class
- scales poorly with class numbers
- performs same as OvA
- DNN more efficient

- **Single NN with multiple outputs**

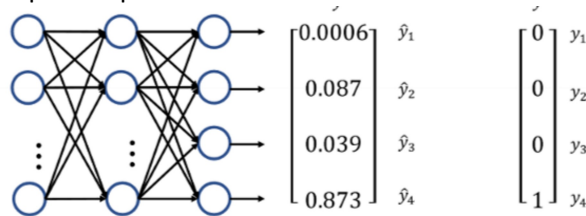
- one NN, change output layer to have 1 node per class
- each output acts as binary classifier for that class (0 or 1)



- can possibly use sigmoid as activation on output layer? interpret as probability

- **Softmax Activation Function**

- normalizes outputs st output nodes produce value bw 0-1 and sum to 1
- can predict probabilities for each class



$$L(\hat{y}, y) = - \sum_{j=1}^{n_c} y_j \log(\hat{y}_j) = ???$$

- produces vector length n_c

$$g_i(Z) = \frac{e^{z_i}}{\sum_{j=1}^{n_c} e^{z_j}}$$

- Categorical Cross Entropy Loss (Softmax Loss) is a generalization of Binary Cross Entropy Loss

$$L(\hat{y}, y) = - \sum_{j=1}^{n_c} y_j \log(\hat{y}_j)$$

- generalization of sigmoid
- Cost function

$$J(W, B) = \frac{1}{m} \sum_{i=1}^m L(\hat{y}^{(i)}, y^{(i)})$$

- Backpropagation

8. Deep Neural Networks

- bias shifts the function, else only lines passing through origin
- 3 layer NN?
- Fully connected FC or dense layers: each input connects to each node
 - each FC layer can have diff number of units
 - diff number of weights&biases?
 - ex. 2 hidden layers:

$$Y(X) = \sigma(W^{[3]} \tanh(W^{[2]} \tanh(W^{[1]}X + B^{[1]}) + B^{[2]}) + B^{[3]})$$

- Feature engineering
 - with LR, can manually transform features to encode non-linearity
 - data is now linearly separable
- Works best in problems with unstructured data as input, or complex relationships but clear goals
- **Backpropagation through softmax and categorical cross entropy loss**