

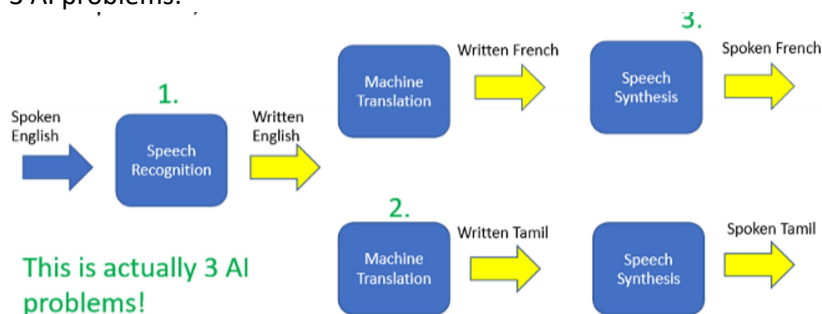
RNN and Natural Language Processing

Tuesday, December 14, 2021 6:52 PM

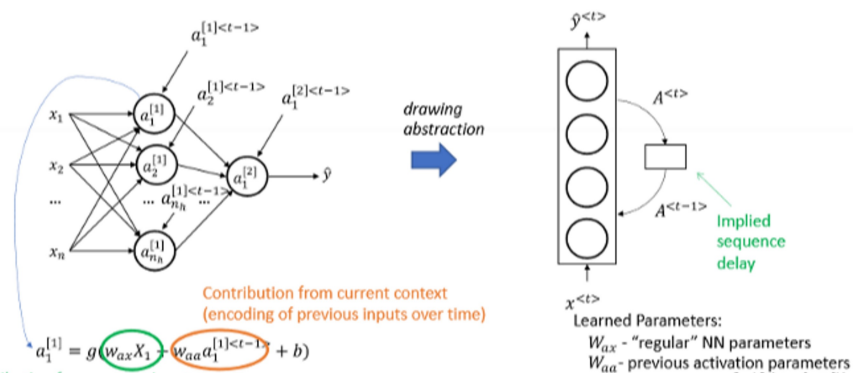
21. Recurrent Neural Networks

Translation

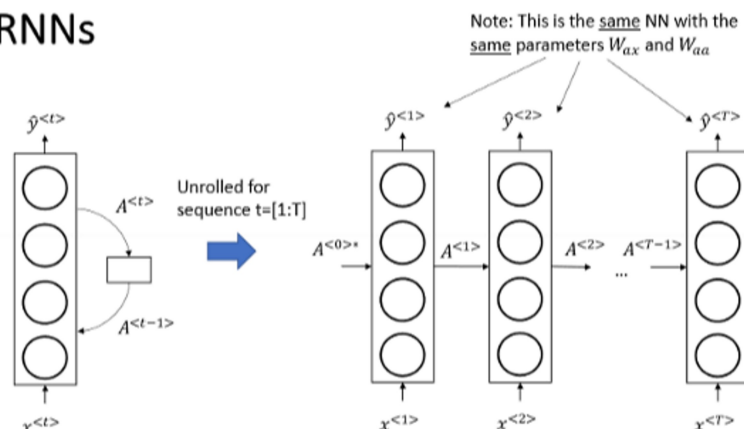
- 3 AI problems:





- information in individual components of data and their ordering
- need to consider the context!
 - o eg. Michael Jordan made a basket
- **Giving ML Context**
 - o understand new info
 - o increasing inputs to system to reflect context (data from the past) doesn't scale!
 - o activations from previous step in sequence can be used to "bias" activations in next step
 - could simultaneously learn amount of context required while we learn input to output mappings




RNNs



- **I/O Sequence Length Flexibility in RNNs**
 - o Flexibility around size of input/output data



- RNN Feature Extraction



- eg. neuron learned to activate strongly to patterns that make an 11x11 swatch of facial hair
- RNNs isolate elements of sequences like conv filters isolate regions of an image

- Bidirectional RNN

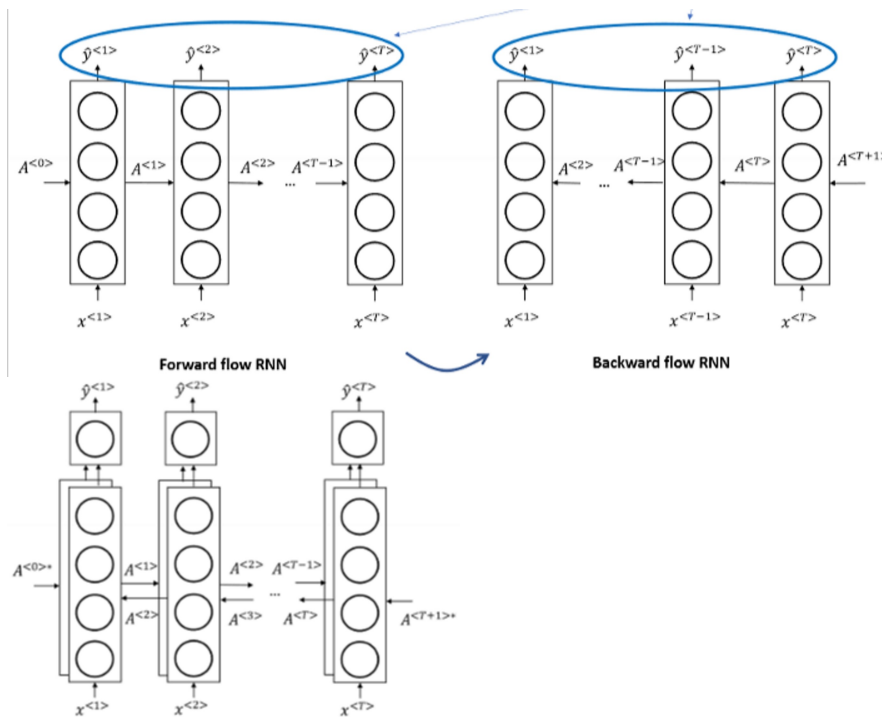
- can look forward and backward in time
- buffer the inputs long enough to consider context in 2 directions
 - context may only be defined later in the sentence

"The fair approach is to split the profits 50/50."

"The fair was in town for only one night."

The context is defined later in the sentence.

- create forward flow RNN and backward flow RNN then combine outputs
- using BRNNs for each sentence = state of the art for NLP



23. Implementation of RNNs

RNN Notation

- Inputs: $x^{(i)<t>}$ where $i = 1:m$ and $t = 1:T_x^{(i)}$ Now, each input and output in the training example has a sequence length T .
- Outputs: $y^{(i)<t>}$ where $i = 1:m$ and $t = 1:T_y^{(i)}$ Like before, we have m training examples.

- eg. for x : "cats are nice", want to make each word an element of our sequence


$$x^{<1>} = \text{"Cats"}, x^{<2>} = \text{"are"}, x^{<3>} = \text{"nice"}, x^{<4>} = \text{"."}$$

- need to assign each word a number
- can create an ordered dictionary, assign each word a num based on position in sequence
 - o may add unintentional bias, words are biased based on position in alphabet

Word	Position
a	1
aaron	2
...	
cats	520
...	
nice	3024
...	
zulu	19,999
<eos>, " ", "	20,000


"Cats": 520
"Are": 97
"Nice": 3024
".": 20,000

Why should "a" and "aaron" have a similar magnitude, but not "a" and "erin", or "tom"? Words are biased together based on their position in the alphabet!



What set of parameters will be able to create a distinct activation for each word?

- need normalized representation and less compressed
 - o instead use one-hot encoding, no order bias




- Unknown words

- one more vector element called unknown word ("UKW")
- works as long as all important words NP task uses are included in vocabulary

RNN Loss Function

- define loss function to optimize later
- define overall loss to be sum of each element of sequence
- using one-hot encodings, re-use Cross Entropy Loss for each element




$$L(\hat{y}, y) = \sum_{t=1}^{T_y} L^{<t>}(\hat{y}^{<t>}, y^{<t>})$$

And, given we are using one-hot encodings, we can re-use Cross Entropy Loss for each element:

$$\begin{aligned} L^{<t>}(\hat{y}^{<t>}, y^{<t>}) &= -y^{<t>} \log(\hat{y}^{<t>}) \\ &\quad - (1 - y^{<t>}) \log(1 - \hat{y}^{<t>}) \end{aligned}$$


- Backprop in RNNs



Let's imagine we have a $T_y = T_x = 2$

1. calculate \hat{y} using computation graph
2. determine the loss
3. update each parameter (using partial derivative of cost)
4. repeat until $J < \text{target}$

Step 1: Calculate \hat{y} using computation graph.




Step 2: Determine the loss.

$$J = -\frac{1}{m} \left(\sum_{i=1}^m y^{(i)} \log(\hat{y}^{(i)}) + \sum_{i=1}^m (1 - y^{(i)}) \log(1 - \hat{y}^{(i)}) \right)$$

Step 3: Update each parameter (Using the partial derivative of cost).

$$w = w - \alpha \frac{\partial J}{\partial w}; b = b - \alpha \frac{\partial J}{\partial b}$$

- Partial Derivatives on RNN




○ Vanishing gradients in RNNs

- as sequences get long, can be difficult to enable earlier elements to correctly influence later outputs
- influence of early elements of the sequences are "washed out" by repeat multiply accumulates in each iteration
- activation unit for RNN:

$$a_1^{[1]<t>} = g(w_{ax}X_1 + w_{aa}a_1^{[1]<t-1>} + b)$$

- repeated applying same RNN activation units. To bypass current activation, we hold previous value (add memory to RNN)



- Simplified Gated Recurrent Unit (GRU)

- get a value between 0 and 1 based on learned param and standard RNN unit inputs

$$\Gamma_{\mu} = \sigma(w_{\mu x}X_1 + w_{\mu a}a_1^{[1]<t-1>} + b_{\mu})$$

- Use function to decide whether keep previous activation or update

$$a_1^{[1]<t>} = \Gamma_{\mu} * \tilde{a}_1^{[1]<t>} + (1 - \Gamma_{\mu}) * a_1^{[1]<t-1>}$$

- standard activation calculation now a candidate:

$$\tilde{a}_1^{[1]<t>} = g(w_{ax}X_1 + w_{aa}a_1^{[1]<t-1>} + b)$$

- similar to creating a mux function

- Long Short Term Memory (LSTM)

- maintain influence of value across many sequences
- manage vanishing gradient problem in RNN
- create new var c, expression of internal memory
- gating: update, forget, output

$$\Gamma_{\mu} = \sigma(w_{\mu x}X_1 + w_{\mu a}a_1^{[1]<t-1>} + b_{\mu}) \quad \text{Update}$$

$$\Gamma_f = \sigma(w_{fx}X_1 + w_{fa}a_1^{[1]<t-1>} + b_f) \quad \text{Forget}$$

$$\Gamma_o = \sigma(w_{ox}X_1 + w_{oa}a_1^{[1]<t-1>} + b_o) \quad \text{Output}$$

- use funcs to manage internal memory and output vals
- eg. create candidate memory:


$$\tilde{c}_1^{[1]<t>} = g(w_{ax}X_1 + w_{aa}a_1^{[1]<t-1>} + b)$$

- update internal memory, update and forget

$$c_1^{[1]<t>} = \Gamma_\mu * \tilde{c}_1^{[1]<t>} + \Gamma_f * c_1^{[1]<t-1>}$$

- create output separate from internal memory

$$a_1^{[1]<t>} = \Gamma_o * \tanh(c_1^{[1]<t>})$$



- structure of many sequential data sets, key element critical for period of time, then no longer relevant

24. Natural Language Processing (NLP)