

Chapter 5 Network Layer: Control Plane

Sunday, April 25, 2021 2:14 PM

5.1 Intro

- how forwarding and flow tables are computed, maintained and installed
 - **Per-router control:** routing algorithm runs in each router, fwd and routing func contained in each router

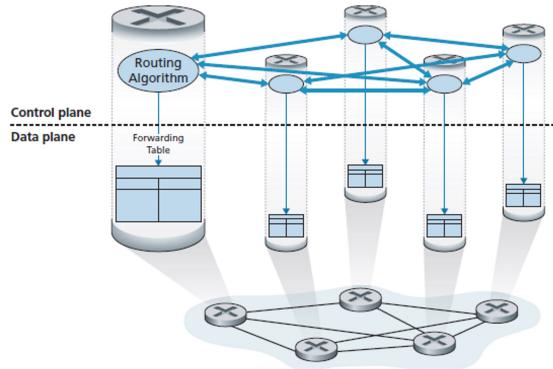


Figure 5.1 • Per-router control: Individual routing algorithm components interact in the control plane

- **Logically centralized control:** logically centralized controller computes and distributes the fwd tables to be used by each router

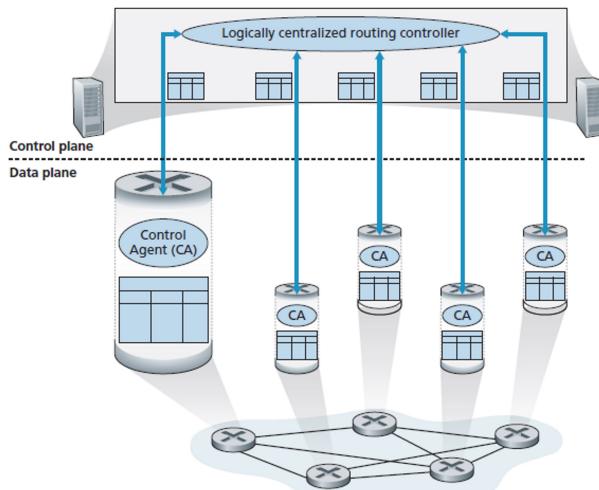
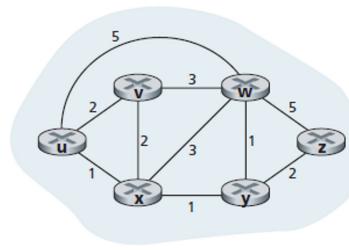


Figure 5.2 • Logically centralized control: A distinct, typically remote, controller interacts with local control agents (CAs)



Abstract graph model of a computer network

5.2 Routing Algorithms

- **Routing protocols:**
 - determine good paths/routes from senders to receivers through network for routers
 - has least cost
- Use graph for routing problems, $G = (N, E)$
 - N = set of routers $\{u, v, w, x, y, z\}$
 - E = set of links $\{(u,v), (u,x), (v,x), (v,w), (x,w), (x,y), (w,y), (w,z), (y,z)\}$
- **Classification**
 - **Centralized/Global**
 - All routers have complete topology, link cost info
 - **Decentralized**
 - iterative process of computation, exchange info with neighbors
 - routers initially only know link costs to attached neighbors
 - distance vector algorithms
 - **Static**
 - routes change slowly over time
 - **Dynamic**

- routes change more quickly, periodic updates or in response to link cost changes

5.2.1 Link-State (LS) Routing Algorithm

- Dijkstra's algorithm

- computes least-cost path from one node (source u) to all other nodes in network
- iterative, after k iterations, know least cost path to k destinations
- **Notation**
 - $D(v)$: cost of least cost path from source node u to dest v during iteration
 - $p(v)$: previous node along current least cost path from source to v
 - N' : subset of nodes. v is in N' if least cost path from source to v is known

1 *Initialization:*

```

2    $N' = \{u\}$            /* compute least cost path from u to all other nodes */
3   for all nodes v
4     if v adjacent to u      /* u initially knows direct-path-cost only to direct neighbors */
5       then  $D(v) = c_{u,v}$     /* but may not be minimum cost */
6     else  $D(v) = \infty$ 
7

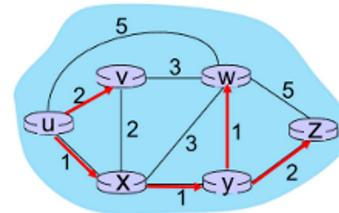
```

8 Loop

```

9   find w not in  $N'$  such that  $D(w)$  is a minimum
10  add w to  $N'$ 
11  update  $D(v)$  for all v adjacent to w and not in  $N'$ :
12     $D(v) = \min(D(v), D(w) + c_{w,v})$ 
13  /* new least-path-cost to v is either old least-cost-path to v or known
14  least-cost-path to w plus direct-cost from w to v */
15 until all nodes in  $N'$ 

```

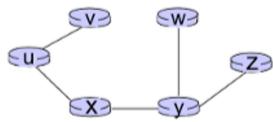


step	N'	$D(v), p(v)$	$D(w), p(w)$	$D(x), p(x)$	$D(y), p(y)$	$D(z), p(z)$
0	u	2, u	5, u	1, u	∞	∞
1	ux	2, u	4, x		2, x	∞
2	uxy	2, u	3, y			4, y
3	uxyv		3, y			4, y
4	uxyw					4, y
5	uxywz					

Initialization (step 0): For all a: if a adjacent to u then $D(a) = c_{u,a}$
 find a not in N' such that $D(a)$ is a minimum
 add a to N'
 update $D(b)$ for all b adjacent to a and not in N' :
 $D(b) = \min(D(b), D(a) + c_{a,b})$

Table 5.1 ♦ Running the link-state algorithm on the network in Figure 5.3

resulting least-cost-path tree from u:



resulting forwarding table in u:

destination	outgoing link
v	(u,v)
x	(u,x)
y	(u,x)
w	(u,x)
x	(u,x)

route from u to v directly

route from u to all other destinations via x

- 0 step initialization

- least cost paths (direct one hop) from u to neighbors v,x,w are 2,1,5
- y and z are inf since not directly connected
- **least cost is x at 1**, add to N' set
- Line 12 of LS algorithm update $D(v)$

- 1 step

- cost of path to v unchanged (least cost is still 2 from u, rather than 3 u,x)
- cost of path to w is 4 (was 5, but least cost is 4 through x)
- cost of path to y is 2
- **least cost is y at 2**, add to N' set

- 2 step

- nodes v and y have least cost path, break tie arbitrarily, add y to set
- cost to remaining nodes not yet in N' are v,w,z

- Complexity

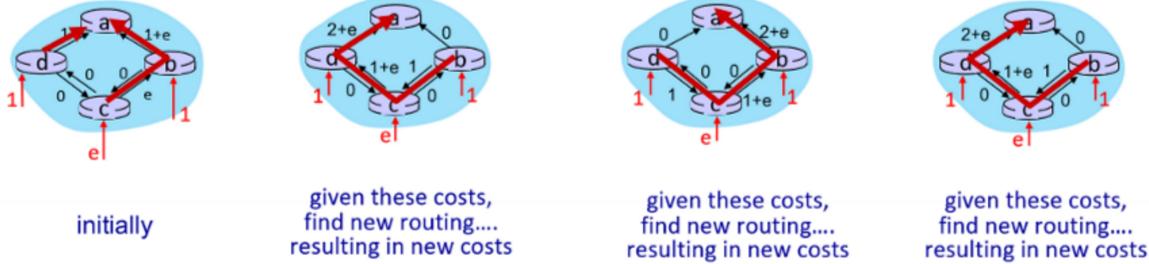
- **Algorithm:** n nodes, $O(n^2)$ complexity, with $n(n+1)/2$ comparisons
- **Message:** each router must broadcast link info to other n routers, crosses $O(n)$ links, overall msg complexity = $O(n^2)$

- Oscillations

- when link costs depend on traffic volume, route oscillations possible
- routing to dest a, traffic entering at d,c,e rates 1, e(<1), 1

- Prevent Oscillations:

- mandate that link costs not depend on amount of traffic carried (unacceptable since need to avoid congestion links)
- ensure that not all routers run the LS algorithm at same time



5.2.2 The Distance-Vector (DV) Routing Algorithm

- iterative async and distributed
 - each node receives some info from one or 2 of neighbors, calculates, then distributes results
 - process continues until no more info exchanged
 - does not require all of nodes to operate in lockstep with each other
- Least cost $d_x(y)$ related by Bellman Ford equation

Bellman-Ford equation

Let $D_x(y)$: cost of least-cost path from x to y .
 Then:

$$D_x(y) = \min_v \{ c_{x,v} + D_v(y) \}$$

v's estimated least-cost-path cost to y

min taken over all neighbors v of x

direct cost of link from x to v

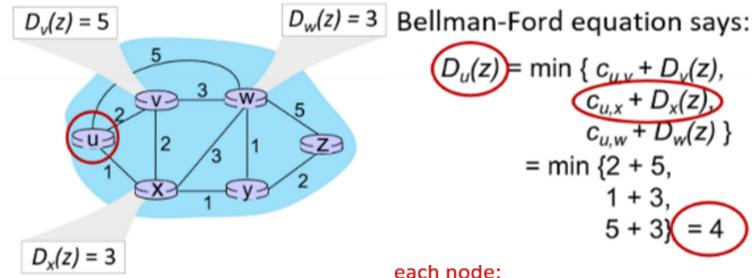
- provides entries in node x 's forwarding table
- node achieving min (x) is next hop on estimated least cost path to dest (z)
- each node sends its own distance vector estimate to neighbors
- when x receives new DV estimate from neighbors, updates its own DV using BF eqn

```

1 Initialization:
2   for all destinations  $y$  in  $N$ :
3      $D_x(y) = c(x,y)$  /* if  $y$  is not a neighbor then  $c(x,y) = \infty$  */
4   for each neighbor  $w$ 
5      $D_w(y) = ?$  for all destinations  $y$  in  $N$ 
6   for each neighbor  $w$ 
7     send distance vector  $D_x = [D_x(y) : y \in N]$  to  $w$ 
8
9 loop
10  wait (until I see a link cost change to some neighbor  $w$  or
11    until I receive a distance vector from some neighbor  $w$ )
12
13  for each  $y$  in  $N$ :
14     $D_x(y) = \min_v (c(x,v) + D_v(y))$ 
15
16 if  $D_x(y)$  changed for any destination  $y$ 
17   send distance vector  $D_x = [D_x(y) : y \in N]$  to all neighbors
18
19 forever
  
```

- LS vs DV algorithms

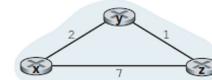
- message complexity
 - LS: n routers, $O(n^2)$ msgs sent
 - DV: exchange bw neighbors, convergence varies
- speed of convergence
 - LS: $O(n^2)$ algorithm, $O(n^2)$ msg, may have oscillations
 - DV: convergence time varies, may have routing loops
- robustness
 - LS: router can advertise incorrect link cost, router computes own table
 - DV: can advertise incorrect path cost - block below



each node:

```

  wait for (change in local link
  cost or msg from neighbor)
  recomput DV estimates using
  DV received from neighbor
  if DV to any destination has
  changed, notify neighbors
  
```



Node x table											
from	x	y	z	cost to	x	y	z	cost to	x	y	z
x	0	2	7	from	0	2	3	from	0	2	3
y	∞	∞	∞	from	2	0	1	from	2	0	1
z	∞	∞	∞	from	7	1	0	from	3	1	0

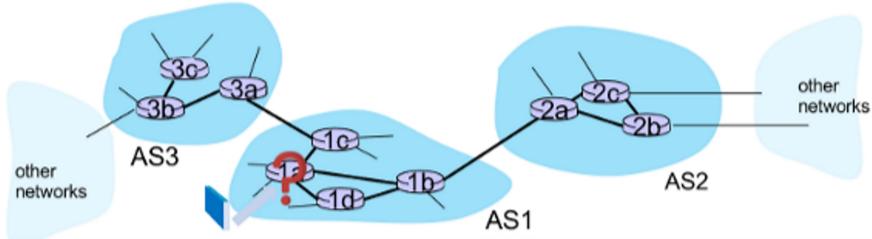
Node y table											
from	x	y	z	cost to	x	y	z	cost to	x	y	z
x	0	2	7	from	0	2	7	from	0	2	3
y	∞	∞	∞	from	2	0	1	from	2	0	1
z	∞	∞	∞	from	7	1	0	from	3	1	0

Node z table											
from	x	y	z	cost to	x	y	z	cost to	x	y	z
x	0	2	7	from	0	2	7	from	0	2	3
y	∞	∞	∞	from	2	0	1	from	2	0	1
z	∞	∞	∞	from	7	1	0	from	3	1	0

- DV: can advertise incorrect path cost - black holing, each routers table used by others, error propagates

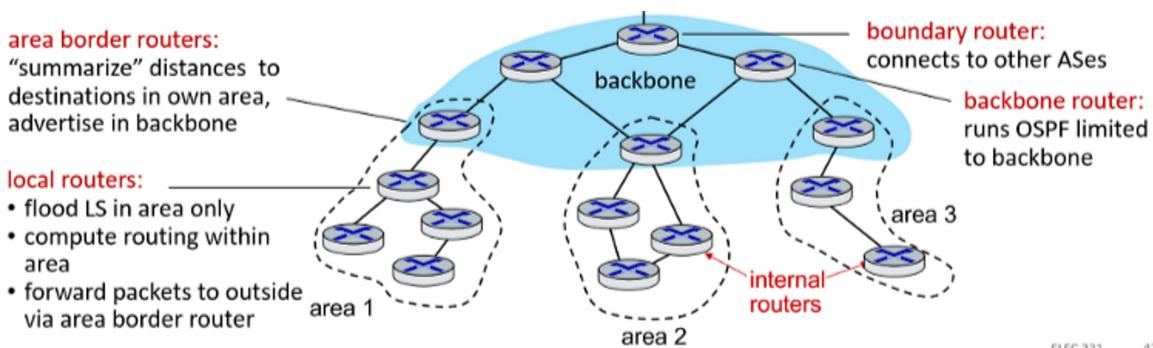
5.2 Intra-AS Routing in the Internet: OSPF

- aggregate routers into regions - Autonomous systems, domains
- **Scale**
 - all routers identical, network flat, but not true in practice
 - cant store all dest in routing tables, exchanges would swamp links
- **administrative autonomy**
 - each network admin may want to control routing in its own network
- **intra-AS**
 - routing among withing same AS
 - all routers must run same intra domain protocol
 - gateway router: at edge of its own AS, has links to routers in other AS
- **inter-AS**
 - routing among AS
 - gateway routers perform inter domain routing
- **Interconnected AS**
 - fwd table configured by intra and inter AS routing algorithms
 - intra AS determine entries for dest within AS
 - inter AS and intra AS determine entries for external dest



- Open Shortest Path First (OSPF)

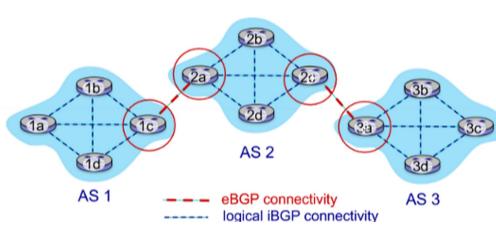
- classic link state routing
 - each router floods OSPF link state advertisements
 - each router has full topology, uses Dijkstra's algorithm to compute fwd table
 - all msgs are authenticated
- two-level hierarchy
 - link state advertisements flooded
 - each node has detailed area topology, only knows direction to reach other dest



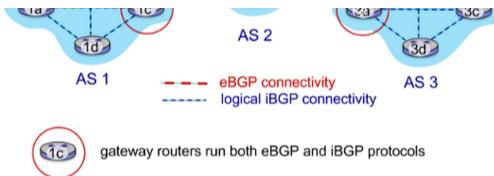
ELEC 331 47

5.4 Routing Among the ISPs: BGP

- allows subnet to advertise its existance, and destination it can reach
- BGP Border Gateway protocol
- Provides each AS a means to
 - eBGP: obtain subnet reachability information from neighboring Ases
 - iBGP: propagate reachability info to all AS internal routers
 - determine good routes

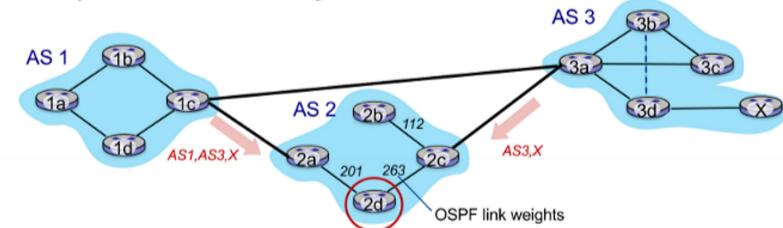


- eBGP: obtain subnet reachability information from neighboring Ases
- iBGP: propagate reachability info to all AS internal routers
- determine good routes



- Hot potato routing

- choose local gateway with least intra domain cost
 - ex. 2d chooses 2a even though more AS hops to X
- tries to reduce cost in its own AS while ignoring other components of the end-to-end costs outside its AS



5.4.5 Routing Policy

- when a router selects a route to a destination, AS routing policy can trump all other considerations, such as shortest AS path or hot potato

- BGP route selection

- router may learn about more than 1 route to destination
- selects based on:
 - local preference value attribute: policy decision
 - shortest AS-PATH
 - closest NEXT-HOP router: hot potato routing
 - additional criteria

5.5 SDN Control Plane